

Lecture 7: Other Applications in Finance and Future Trends

Large Language Models in Finance

Juan F. Imbet

Paris Dauphine – PSL University

Outline

Where We Have Been: A Map of the Book

Ch.	Topic	Primary Technique	Application
1	Text Fundamentals	TF-IDF, word embeddings	Sentiment, classification
2	Modern LLM Landscape	Transformer, RLHF	General financial NLP
3	Training & Fine-Tuning	LoRA, PEFT	Domain adaptation
4	LLM Agents	ReAct, tool use	Research automation
5	Business Valuation	RAG, summarisation	DCF, comparables
6	Credit Risk	Constrained gen., SHAP	Scoring, fairness
7	Applications & Future	Workflows, multimodal	Deployment, governance

The through-line

Early chapters: lightweight representations, large corpora, modest compute.

Later chapters: generative and agent-based models, multi-step reasoning.

This chapter: cross-cutting concerns of **deployment, governance, and frontier capability**.

Remaining Application Areas

Portfolio Management

- Natural-language mandate specification → quantitative optimisation inputs
- ESG screening from heterogeneous sustainability reports
- Earnings surprise prediction from call transcripts: outperforms bag-of-words dictionaries
- López-Lira (2025): heterogeneous LLM agents in simulated markets → real market dynamics + systemic risk from correlated strategies

Regulatory Compliance

- MiFID II suitability report generation (client profile + instrument + regulatory template)
- AML Suspicious Activity Reports: improved consistency, reduced analyst time

Insurance: underwriting risk extraction; claims triage; policy question-answering.

Algorithmic Trading: signal generation (news/filings → sentiment)

Choosing the Right Deployment Pattern

Pattern	Best for	Data	Latency	Cost
Zero-shot	General reasoning, novel tasks	None	Low	Low
RAG	Document QA, live data, audit trail	Corpus	Medium	Medium
Fine-tuning	Structured output, consistency	Labelled	Low	High

Three practical heuristics:

- 1 Prefer RAG over fine-tuning when the task requires grounding in *updatable* documents (current regulation, recent filings). Fine-tuning embeds knowledge statically.
- 2 Fine-tune when output format is highly structured and consistency across thousands of calls matters more than generality.
- 3 Zero-shot with a frontier model is a valid baseline before investing in more expensive approaches. FinanceBench: frontier models answer $\approx 80\%$ of financial QA correctly on the first attempt.

Architecture Family Selection

Task Property	Enc-only	Dec-only	Enc-Dec	Age
Labelled fine-tuning data available	Strong	Weak	Strong	We
Output is classification / embedding	Strong	Weak	Weak	We
Output requires fluent generation	Weak	Strong	Strong	Str
Task needs up-to-date external data	Weak	Weak	Weak	Str
Real-time / low-latency constraint	Strong	Weak	Weak	We
Audit trail / interpretability needed	Strong	Weak	Weak	We
Task spans multiple steps / documents	Weak	Weak	Weak	Str

Archetypes:

- Trade surveillance classifier → encoder-only + fine-tuning
- Regulatory filing summariser → decoder-only or enc-dec + RAG
- Autonomous due-diligence agent → agentic + zero/few-shot

Financial Workflow as a Directed Acyclic Graph

Definition

A *financial workflow* is a DAG $W = (T, E)$ where:

- $T = \{t_1, \dots, t_n\}$ is a set of tasks, each $t_i : \mathcal{X}_i \rightarrow \mathcal{Y}_i$
- $E \subseteq T \times T$: $(t_i, t_j) \in E$ iff t_j requires an output of t_i
- Acyclic: no directed path from any task back to itself

An *execution* is a topological ordering consistent with E .

Three canonical automation patterns:

Pattern	Concurrency	Best for
Sequential pipeline	None	Ordered, batch, cascading failures
Event-triggered	Per-event	Semantic triggers, alert systems
Orchestrator-worker	Parallel subtasks	Expert specialisation, replanning

Three Representative Finance Workflows

1. **Earnings call processing pipeline** (sequential):

- 1 Retrieve transcript
- 2 Extract structured guidance (revenue, margin, capex) as JSON
- 3 Compare against analyst consensus; compute surprises
- 4 Classify management tone (constructive / cautious / evasive)
- 5 Draft standardised analyst briefing

Minutes from call end to delivered briefing, vs. hours of analyst time.

2. Regulatory filing monitor (event-triggered): SEC EDGAR RSS feed → classify 8-K by type and materiality → route to appropriate team with summary. Materiality is semantic, not syntactic.

3. Trade surveillance narrative (orchestrator-worker): Rule-based system flags unusual order pattern → LLM retrieves news + social media + internal communications → drafts structured narrative → compliance analyst reviews and decides.

Integration principle

LLMs operate at seconds-to-minutes scale. Execution systems operate at

Event-Triggered EDGAR Monitor: Code Sketch

```
CLASSIFICATION_PROMPT = """Classify this newly filed SEC 8-K.  
Company: {company}  
Filing description: {description}  
  
Respond with valid JSON:  
{"materiality_score": 1-5, "event_type": "earnings|  
    merger_acquisition|...",  
    "summary": "...}"""  
  
def classify_filing(client, company, description):  
    message = client.messages.create(  
        model="claude-opus-4-5",  
        max_tokens=256,  
        messages=[{"role": "user", "content":  
                    CLASSIFICATION_PROMPT.format(  
                        company=company, description=description)}]  
    )  
    return json.loads(message.content[0].text)  
  
def monitor_edgar(poll_interval=300):  
    seen_ids = set()  
    while True:  
        feed = feedparser.parse(EDGAR_RSS)  
        for entry in feed.entries:
```

Definition

$\mathcal{A} = (\mathcal{R}, \mathcal{M}, \mathcal{K}, \mathcal{T}, \mathcal{I})$ where:

- $\mathcal{R}$: retrieval engine over corpus $\mathcal{D}$
- $\mathcal{M}$: LLM reader synthesising response from query + retrieved passages
- $\mathcal{K}$: persistent memory (user preferences, portfolio context)
- $\mathcal{T}$: tool layer (market data APIs, filing databases, calculators)
- $\mathcal{I}$: interface layer (text, voice, tables, visualisations)

Hybrid retrieval is non-negotiable.

Semantic search alone fails on:

- Numerical queries (“what was Apple’s FCF in fiscal 2023?”)
- Exact-term queries (“Article 22 GDPR”, “IFRS 9 ECL”)

Reciprocal Rank Fusion:

$$\text{RRF}(q, d_i) = \frac{1}{k} + \frac{1}{k} \dots$$

RAG Implementation for Financial Documents

Chunking strategy matters. Naive 512-token chunking:

- Fragments financial tables (destroys row/column structure)
- Splits financial statements across chunks
- Disrupts cross-references between filing sections

Section-aware chunking for 10-K filings:

- Each SEC Item indexed independently (Item 1A, 7, 8, etc.)
- Tables extracted as structured objects; indexed by a separate retriever
- Substantial reduction in retrieval noise on financial document QA benchmarks

Source attribution is a non-negotiable requirement:

- Every claim must cite a retrieved passage
- Claims without retrieved support must be flagged as model inference, not retrieved fact
- In regulated contexts: source attribution = the audit mechanism

Context length paradox (Balogh & Didisheim, 2025)

OpenClaw: Local Personal AI for Finance Professionals

Design principles:

- ➊ **Privacy by default:** all processing local; financial data never leaves the machine
- ➋ **Meet users where they work:** queries via WhatsApp, Telegram, Slack, Signal
- ➌ **Model flexibility:** cloud (Claude, GPT-4) for capability; local (Llama, Mistral) for sensitivity
- ➍ **Extensibility through skills:** a Python function with a descriptive docstring becomes a callable tool

Finance-specific skills:

Skill	What it does
/earnings AAPL	Guidance summary vs. consensus from most recent transcript
Filing monitor	Daily digest of new SEC filings from watchlist
Portfolio Q&A	Natural-language questions about local CSV portfolio
Macro digest	Morning briefing from central bank speeches + market moves

Sources of Bias in LLM-Based Finance Systems

- 1 **Training data bias:** frontier LLMs are trained on English-language, US-centric financial content. Geographic bias in risk assessment is a *systematic pricing error*, not merely a fairness concern.
- 2 **Historical discrimination encoded in data:** LLMs trained on historically discriminatory lending records reproduce those patterns. This is not a model failure — it faithfully reproduces the signal in the training data.
- 3 **Feedback loop bias:** when LLM outputs influence asset prices, those prices enter future training data, reinforcing the model's priors. Analogous to Soros's reflexivity but at faster timescales. Widespread deployment of similar strategies amplifies volatility.
- 4 **Anchoring bias:** LLMs weight salient numerical anchors disproportionately. An LLM that just processed a dramatic earnings beat will over-weight that number in subsequent analysis. Documented in human judgment by Tversky and Kahneman (1974); reproduced by LLMs trained on human-generated text.

Fairness Metrics and the Impossibility Theorem

Demographic parity: $\Pr(\hat{Y} = 1 \mid A = a) = \Pr(\hat{Y} = 1 \mid A = b)$ for all a, b .

Equalized odds: $\Pr(\hat{Y} = 1 \mid Y = y, A = a) = \Pr(\hat{Y} = 1 \mid Y = y, A = b)$ for $y \in \{0, 1\}$.

Calibration within group: $\Pr(Y = 1 \mid \hat{P} = p, A = a) = p$ for all p .

Chouldechova (2017) Impossibility Theorem

When base rates differ across groups —

$\Pr(Y = 1 \mid A = a) \neq \Pr(Y = 1 \mid A = b)$ — a non-trivial classifier **cannot simultaneously** satisfy demographic parity, equalized odds, and calibration within group.

Demographic parity and calibration within group are mutually incompatible whenever base rates differ.

Implication

Fairness is a *family of criteria*, not a single concept. Choosing which criterion to satisfy is a policy choice, not a technical one. In regulated

Bias Mitigation: Three Stages

Pre-processing:

- Data rebalancing (oversampling underrepresented groups)
- Causal data augmentation: vary protected attributes while holding other features fixed
- Effect: reduces representation bias; some cost to majority-group accuracy

In-processing:

- Adversarial debiasing: classifier trained to predict outcome while adversary cannot recover sensitive attribute from intermediate representations
- Requires sensitive attribute annotations in training data

Post-processing:

- Threshold adjustment: select decision thresholds satisfying the desired fairness criterion on a held-out calibration set
- Applicable to black-box APIs; comes at the cost of accepting different error rates for different groups

The Regulatory Landscape for AI in Finance

Framework	Key implications for LLM finance
EU AI Act (2024)	Credit scoring and life insurance risk assessment <i>high-risk</i> AI. Conformity assessment before market entry, mandatory registration, transparency obligation, meaningful human oversight.
ECOA / Fair Housing Act	Adverse action notice requirements: specific and accurate, not generic. In direct tension with LLM opacity.
GDPR Art. 22	Right to human review of automated decisions. “Use a large language model” does not constitute meaningful information about the logic involved.
MiFID II	LLM-generated investment recommendations and market analyses must be logged: prompt, model version, output, timestamp.
FSB (2024)	AI amplifies third-party concentration risk, market correlations via herding, and financial disinformation. Call for enhanced oversight.

SR 11-7 for LLM Systems: Three Lines of Defence

SR 11-7 (2011): when an LLM produces quantitative estimates feeding into material business decisions, it constitutes a model — regardless of its internal architecture.

- ① **First line (Development):** document purpose, scope, data inputs, architecture (foundation model + fine-tuning), known limitations, performance metrics. For LLMs: document failure modes from red-teaming explicitly, even if incomplete.
- ② **Second line (Independent Validation):** adversarial testing including prompt injection, jailbreaking, and distribution shift tests. The FAIR framework (Noguer, 2025) addresses LLM-specific failure modes: temporal data inconsistencies, hallucination in numerical extraction, privacy leakage.
- ③ **Third line (Internal Audit):** periodically verify that documentation is current, validation is genuinely independent, findings are tracked to resolution, and model usage has not drifted beyond the validated scope.

Explainability, Monitoring, and Incident Response

Two kinds of explainability:

- **Post-hoc attribution** (SHAP, attention saliency): “what features most influenced this output?”
- **Procedural transparency** (chain-of-thought): provides an auditable reasoning trace, but is a *natural-language reconstruction* of reasoning, not an explanation of neural network internals

Four monitoring signal classes:

- 1 Input distribution: KL divergence between prompt embeddings detects distributional shift
- 2 Output quality: accuracy on sampled outputs where ground truth is verifiable
- 3 Operational: latency, cost, error rates, refusal rates
- 4 Fairness: periodic re-runs of bias evaluation framework

Incident severity classification (pre-define before incidents occur):

Low Hallucination in low-stakes summarisation → standard review cycle

Multimodal LLMs and Long-Context Models

The limitation of text-only models in finance:

- Earnings presentations contain charts visualising trends
- Annual reports embed tables organising financial statements
- Earnings calls carry prosodic cues (hesitation, urgency) lost in transcription

Multimodal models jointly encode text, images, audio, and structured data. Early benchmark evidence: multimodal models outperform text-only on financial QA tasks requiring chart and table interpretation.

Long-context models (hundreds of thousands of tokens): an entire 10-K can be processed in a single forward pass, linking risk factors to MD&A to legal proceedings without losing context between chunks.

The “lost in the middle” problem

Recall accuracy for long-context models degrades for passages positioned far from either end of the context. Retrieval-augmented approaches do not share this limitation.

Autonomous Financial Agents: Opportunities and Risks

The trajectory:

- Quantitative research: hypothesis → tested strategy, order-of-magnitude faster
- Trading strategy self-improvement: propose modifications to signal construction, back-test, adopt improvements that pass quality thresholds

Systemic risk concerns (FSB 2024, López-Lira 2025):

- Many market participants deploying similar LLM strategies → correlated actions
- Cascade risk: simultaneous repositioning in response to a common signal
- Regulators examining circuit-breaker requirements for AI-driven strategies

Governance principle: three output categories

- 1 **Informational:** analysis and recommendations

Five Open Research Problems

- 1 **Calibrated uncertainty quantification:** well-calibrated uncertainty from token-level log-probabilities (Chen et al., 2024) outperforms post-hoc sampling on financial QA. Essential for risk-adjusted decisions and regulatory documentation.
- 2 **Comparable selection in thin markets:** embedding spaces trained on US financial text are biased toward US market characteristics. Cross-market comparables require market-invariant, firm-sensitive embeddings.
- 3 **Cross-lingual finance LLMs:** current multilingual models perform substantially worse on non-English regulatory filings (J-GAAP, HGB, local GAAP variants). Corpora and evaluation frameworks do not yet exist.
- 4 **Feedback loops and memorisation:** Didisheim & López-Lira (2025) formalise the LLM's "predictable memory" — look-ahead bias most severe at low frequencies for large models. Credible identification strategies for loop dynamics are an open problem.
- 5 **Faithful chain-of-thought verification:** stated reasoning in CoT

The Path Forward: Three Disciplines

Why capability and responsibility are jointly necessary

An LLM credit model that discriminates unlawfully misallocates capital.
A trading system amplifying volatility damages market infrastructure.
A valuation agent that hallucinates financial figures undermines the professional standards that give financial analysis its social function.

Three disciplines that separate successful deployments from failures:

- 1 **Measure everything:** retrieval quality, output faithfulness, fairness metrics, distributional drift. What is not measured cannot be managed, and in regulated finance what is not managed creates liability.
- 2 **Document honestly:** purpose, limitations, known failure conditions, the human oversight structure that catches failures. Documentation is the mechanism by which institutional knowledge survives model updates and team turnover.
- 3 **Escalate appropriately:** build human review into every consequential decision workflow as a genuine check, not a legal formality, that

Course Summary: The LLM Finance Stack

Layer	What we have learned
Text & sentiment	Embeddings, sentiment dictionaries, event extraction
Foundation models	Transformer, attention, pre-training, RLHF, fine-tuning
Domain adaptation	LoRA, PEFT, instruction tuning, training data curation
Agents	ReAct, tool use, orchestrator-worker, multi-agent systems
Business valuation	DCF pipeline, scenario analysis, comparables, Monte Carlo
Credit risk	Constrained generation, calibration, SHAP, fairness, SR 11-7
Workflows	DAG patterns, event triggers, EDGAR integration
Research assistant	Hybrid RAG, persistent memory, OpenClaw
Governance	SR 11-7, EU AI Act, GDPR, MiFID II, incident response
Frontier	Multimodal, long-context, autonomous agents, open problem

The field is young

Researchers and practitioners reading this chapter will shape its norms, its standards, and its failures as much as its successes. Use that privilege deliberately.