

# Practical Session 09: Financial NLP and Sentiment

## Hands-On — LM Scoring, FinBERT vs. Lexicon, and a Filing-Date Event Study

Juan F. Imbet

Paris Dauphine – PSL University

June 20, 2026

# Session Overview

**Duration:** 1 hour (50 min work + 10 min debrief)

## Agenda:

- ➊ **Recap** (5 min): LM net sentiment, macro-F1, CAR.
- ➋ **Problem 1** (12 min): Score an MD&A excerpt with the LM lexicon by hand.
- ➌ **Problem 2** (12 min): Macro-F1 vs. accuracy on an imbalanced 3-class confusion matrix.
- ➍ **Problem 3** (8 min): Build a CAR[1, 60] predictive regression.
- ➎ **Case Study** (15 min): Run the chapter-09 companion notebook — LM lexicon figure + net sentiment.
- ➏ **Discussion** (8 min): Lexicon vs. FinBERT vs. LLM in production.

## Code

`code/notebooks/09-financial-nlp-sentiment/` — `demo.ipynb` (deterministic) and `exercises.ipynb` (live EDGAR lab).

# Recap: The Three Formulas You Need

**LM net sentiment** (Problem 1):

$$S_{\text{LM}}(d) = \frac{N_{\text{pos}}(d) - N_{\text{neg}}(d)}{N_{\text{words}}(d)}$$

**Macro-F1** over 3 classes (Problem 2):

$$F1_c = \frac{2 P_c R_c}{P_c + R_c}, \quad \text{macro-F1} = \frac{1}{3} \sum_c F1_c$$

with  $P_c = \frac{TP_c}{TP_c + FP_c}$ ,  $R_c = \frac{TP_c}{TP_c + FN_c}$ .

**Cumulative abnormal return** (Problem 3):

$$\text{CAR}_i[1, 60] = \sum_{\tau=1}^{60} (R_{i,\tau} - \hat{R}_{i,\tau}), \quad \text{CAR}_i[1, 60] = \alpha + \beta_1 \text{SUE}_{i,0} + \beta_2 \tilde{S}_{i,0} + \varepsilon_i$$

## Recap: Standardize Before You Compare

Raw scores are not comparable across firms or time. Z-score within industry-year:

$$\tilde{S}_{i,t} = \frac{S_{i,t} - \mu_{j(i),t}}{\sigma_{j(i),t}}$$

### Why it matters for the regression

A defence or chemicals firm has structurally negative LM language. Without the  $\tilde{S}$  transform,  $\beta_2$  picks up an industry fixed effect, not a sentiment signal. The event study uses  $\tilde{S}$ , never the raw  $S_{LM}$ .

**Also check:** ADF/KPSS for long series; one fixed model vintage to avoid sentiment inflation; align timestamps to market hours (no look-ahead).

# Problem 1: LM Scoring by Hand

**Excerpt** (10-K MD&A):

*“Operating margins contracted as input costs rose. Management remains uncertain about demand but is confident the restructuring will deliver improved profitability.”*

**Tasks.** Using the LM lists (negative dominates; *uncertain* is on the uncertainty list; cost is **neutral** in finance, not negative):

- 1 Identify LM negative, positive, and uncertainty hits.
- 2 Take  $N_{\text{words}} = 24$ . Compute  $S_{\text{LM}}$ .
- 3 A Harvard-GI classifier would also flag *costs* and *contracted*. Recompute the GI-style score. How large is the distortion?

**Discussion:** which word does the lexicon get *wrong* on negation, and how would FinBERT handle it?

# Solution: Problem 1

## LM hits:

- Negative: *contracted*, *uncertain* (also uncertainty) → count *contracted* as 1 negative.
- Positive: *confident*, *improved* = 2.
- Uncertainty: *uncertain* = 1 (hedging signal, not polarity).
- *costs*, *rose* are **neutral** in LM.

$$S_{LM} = \frac{2 - 1}{24} = +0.042 \quad (\text{mildly positive})$$

**GI-style:** adds *costs* and *contracted* as negative → negatives = 3:

$$S_{GI} = \frac{2 - 3}{24} = -0.042$$

**Sign flips** — the GI false positives on neutral finance vocabulary reverse the conclusion. This is exactly the Loughran–McDonald (2011) critique.

**Negation:** “not confident” would still score *positive* under any lexicon; FinBERT’s contextual [CLS] embedding resolves it.

## Problem 2: Macro-F1 vs. Accuracy

A 3-class classifier on 1,000 financial sentences (neutral-heavy). Confusion matrix (rows = true, cols = predicted):

|          | pred neg | pred neu | pred pos | total |
|----------|----------|----------|----------|-------|
| true neg | 60       | 35       | 5        | 100   |
| true neu | 20       | 560      | 20       | 600   |
| true pos | 5        | 90       | 205      | 300   |

### Tasks.

- 1 Compute overall accuracy.
- 2 Compute per-class precision, recall, F1 for **negative**.
- 3 Given  $F1_{\text{neu}} \approx 0.86$  and  $F1_{\text{pos}} \approx 0.74$ , compute macro-F1.
- 4 Why is the gap between accuracy and macro-F1 the point of this exercise?

## Solution: Problem 2

$$\text{Accuracy} = (60 + 560 + 205)/1000 = \mathbf{0.825}.$$

### Negative class:

- $TP = 60$ ,  $FP = 20 + 5 = 25$ ,  $FN = 35 + 5 = 40$ .
- Precision =  $60/85 = 0.706$ ; Recall =  $60/100 = 0.600$ .
- $F1_{\text{neg}} = \frac{2(0.706)(0.600)}{0.706 + 0.600} = \mathbf{0.649}$ .

$$\text{Macro-F1} = \frac{1}{3}(0.649 + 0.86 + 0.74) = \mathbf{0.750}.$$

### The lesson

Accuracy 0.825 looks strong, but it is propped up by the 600 easy neutral cases. Macro-F1 0.750 reveals the model is markedly weaker on the economically important **negative** class ( $F1 = 0.65$ ). For a trading or risk signal, the neg/pos tails are what you act on — so report macro-F1.



## Problem 3: Designing the Event Study

**Setup.** 800 earnings calls. For each you have: the call-day standardized text sentiment  $\tilde{S}_{i,0}$ , the earnings surprise  $SUE_{i,0} = (EPS - \widehat{EPS})/\hat{\sigma}_i$ , and daily returns.

### Tasks.

- 1 Write  $CAR_i[1, 60]$  in terms of abnormal returns. What benchmark  $\hat{R}$  would you use, and why not raw returns?
- 2 State the regression and the null/alternative on  $\beta_2$ .
- 3 You find  $\hat{\beta}_2 = 0.018$  ( $t = 2.6$ ) after controlling for SUE. Interpret economically.
- 4 Name two robustness checks before you trust this.

## Solution: Problem 3

1.  $CAR_i[1, 60] = \sum_{\tau=1}^{60} (R_{i,\tau} - \hat{R}_{i,\tau})$ ,  $\hat{R}$  from a market model / CAPM / factor benchmark. Raw returns mix in market-wide moves; abnormal returns isolate the firm-specific drift.
2.  $CAR_i[1, 60] = \alpha + \beta_1 SUE_{i,0} + \beta_2 \tilde{S}_{i,0} + \varepsilon_i$ .  $H_0 : \beta_2 = 0$  (text adds nothing beyond the EPS number);  $H_1 : \beta_2 > 0$  (residual text sentiment predicts drift).
3.  $\hat{\beta}_2 = 0.018$ ,  $t = 2.6$ : a one-SD increase in standardized sentiment  $\rightarrow$  +1.8 pp of 60-day abnormal return, *orthogonal* to the quantitative surprise. Text carries incremental information — post-announcement drift conditioned on tone.
4. Robustness: (i) strictly out-of-sample period not used for model selection (economic-validity test); (ii) cluster SEs by firm/date; also check look-ahead (timestamp vs. market close) and a placebo pre-event window.

# Case Study: Chapter-09 Companion Notebook

Open code/notebooks/09-financial-nlp-sentiment/demo.ipynb. It is deterministic (no network) and reproduces the LM lexicon figure + the worked net-sentiment example.

```
import matplotlib.pyplot as plt

# Loughran-McDonald finance dictionary word-list sizes.
lists = {"Negative": 2355, "Litigious": 903, "Positive":
        354,
        "Uncertainty": 297, "Weak modal": 27, "Strong modal": 19}

items = sorted(lists.items(), key=lambda kv: kv[1])
plt.barh([k for k, _ in items], [v for _, v in items], color="#4c72b0")
plt.xlabel("Number of words"); plt.title("Loughran-McDonald dictionary")
plt.tight_layout(); plt.show()

# Net LM sentiment for the chapter M&A excerpt (ex:lm-example).
n_pos, n_neg, n_words = 1, 4, 35
```

# Case Study: Diagnostic Questions

Run `demo.ipynb`, then answer:

**Q1.** The negative list (2,355) dwarfs the positive list (354). What does this asymmetry tell you about the language of corporate disclosure — and about the *base rate* a lexicon will produce on a typical 10-K?

**Q2.** The excerpt scores  $S_{LM} = -0.086$ . Re-run with `n_pos`, `n_neg` = 2, 2. How sensitive is the sign to a single word-list membership decision? What does this imply for reproducibility?

**Q3 (live lab).** Move to `exercises.ipynb`: it pulls real 10-Ks from EDGAR and prices from `yfinance`. Score one filing, then test whether filing-day LM negativity is associated with the filing-day abnormal return. Does the sign match Loughran–McDonald (2011)?

## Extension

Swap the lexicon for `yyianghkust/finbert-tone` on the same filings. Where do the two methods *disagree*, and which sentences drive it?

# Discussion: Choosing a Method in Production

For each scenario, pick lexicon / FinBERT / LLM and justify on cost, latency, labels, interpretability, and accuracy.

- 1 A regulator must audit *why* every sentiment score on a 10-K was assigned. 50k filings/year, no labelled data, full transparency required.
- 2 A desk scores 5M news headlines/day with a  $< 50$  ms/headline budget and 10k labelled examples in hand.
- 3 A research team prototypes sentiment on a brand-new document type (crypto whitepapers) with *zero* labels and a small corpus.

## Frame your answer around

Auditability (lexicon wins), throughput/latency (distilled FinBERT wins), cold-start flexibility (LLM zero-shot wins) — and when you would *distil* an LLM into a small model for scale.

# Wrap-Up: Key Takeaways

- **Lexicons are auditable but brittle:** the GI vs. LM sign flip in Problem 1 is the whole case for domain-specific dictionaries.
- **Report macro-F1, not accuracy:** the neutral majority hides weakness on the actionable tails (Problem 2).
- **Standardize, then test economically:** z-score within industry-year, run the CAR regression, validate strictly out-of-sample (Problem 3).
- **Method choice is a trade-off:** auditability vs. throughput vs. cold-start flexibility — no universal winner.

## Next practical

Practical 10 — domain-specific LLMs: fine-tuning and adaptation of FinBERT's successors on long financial documents.

# APPENDIX

## Stretch Problems

Extra / optional material — beyond the core lecture.

## Stretch: Krippendorff's Alpha

Two annotators label 5 sentences (P=pos, N=neu, G=neg):

|          | s1 | s2 | s3 | s4 | s5 |
|----------|----|----|----|----|----|
| Annot. A | P  | N  | G  | N  | P  |
| Annot. B | P  | N  | G  | G  | P  |

**Tasks.** (1) Compute observed disagreement  $D_o$  (fraction of disagreeing pairs). (2) With marginal frequencies, compute  $D_e = 1 - \sum_c (n_c/n)^2$ . (3) Compute  $\alpha = 1 - D_o/D_e$ . Is it above Krippendorff's 0.667 reliability bar? (4) Why would the *ordinal* variant give a different, more lenient verdict on the s4 P/G... here N/G disagreement?



# Stretch: From Score to Sharpe

**Setup.** A standardized news-sentiment signal  $\tilde{S}$  is rebalanced daily into a long–short portfolio (long top decile, short bottom decile).

## Tasks.

- 1 The in-sample annualized Sharpe is 1.4; strictly out-of-sample it is 0.5. Name three reasons for the decay (hint: one is in the lecture’s “sentiment inflation” box).
- 2 Kirtac & Germano (2024) report OPT beats the LM dictionary on long–short Sharpe. Why might the gap shrink *after* transaction costs at high turnover?
- 3 Lehner (2024) shows LLMs can rewrite filings to shift measured tone. How does adversarial text-generation threaten a filings-based signal’s out-of-sample stability?